Advanced R

September 2016

Hadley Wickham
@hadleywickham
Chief Scientist, **RStudio**

HELLO

my name is

Hadley

Sean Hackett

Postdoc,
Storey Lab,
Princeton

Jared Lander

Chief Data Scientist,
Lander Analytics

Your turn

This means that you have to work!

This course is very hands on, and while we're here to help you, the best resource is often the person sitting next to you.

Introduce yourself to your neighbours. Who are you and what are you using R for?

Today you'll learn

R foundations

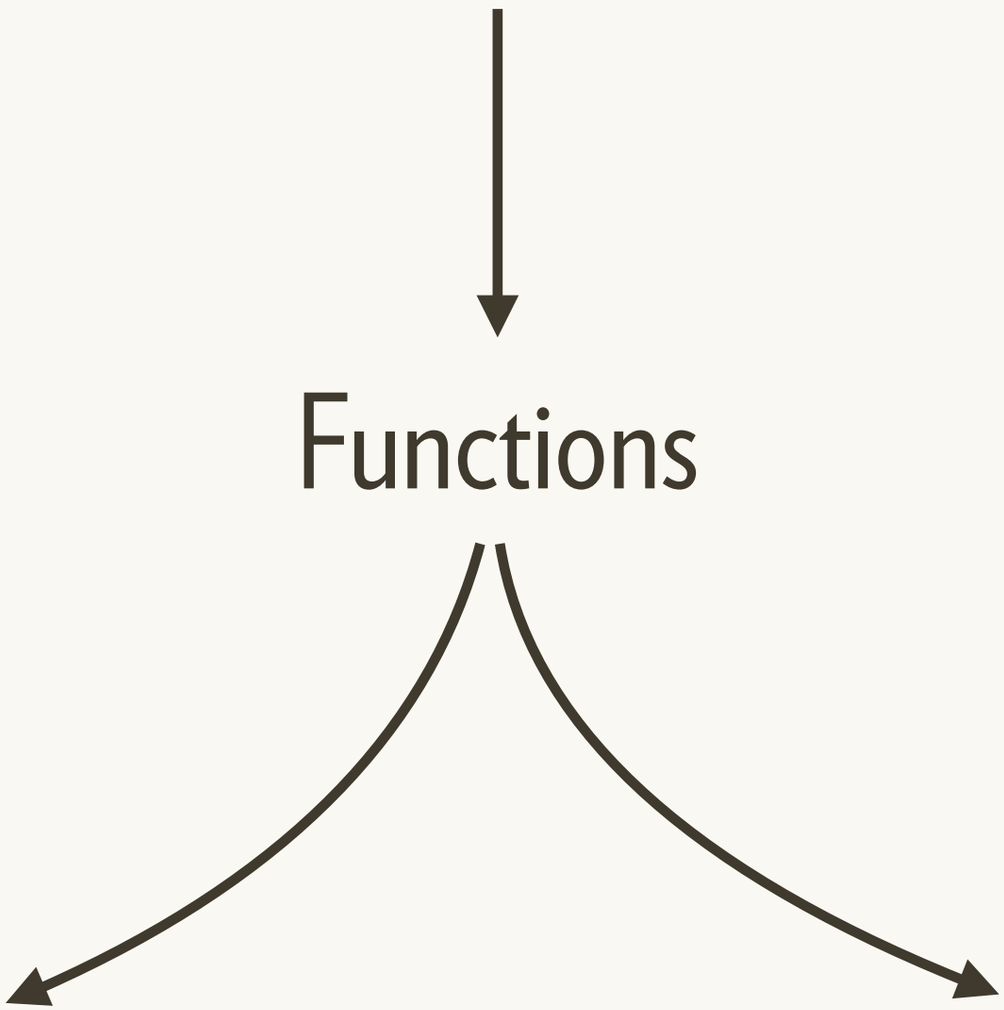

Functions

Functional programming

Object oriented programming

The material is mostly drawn from:

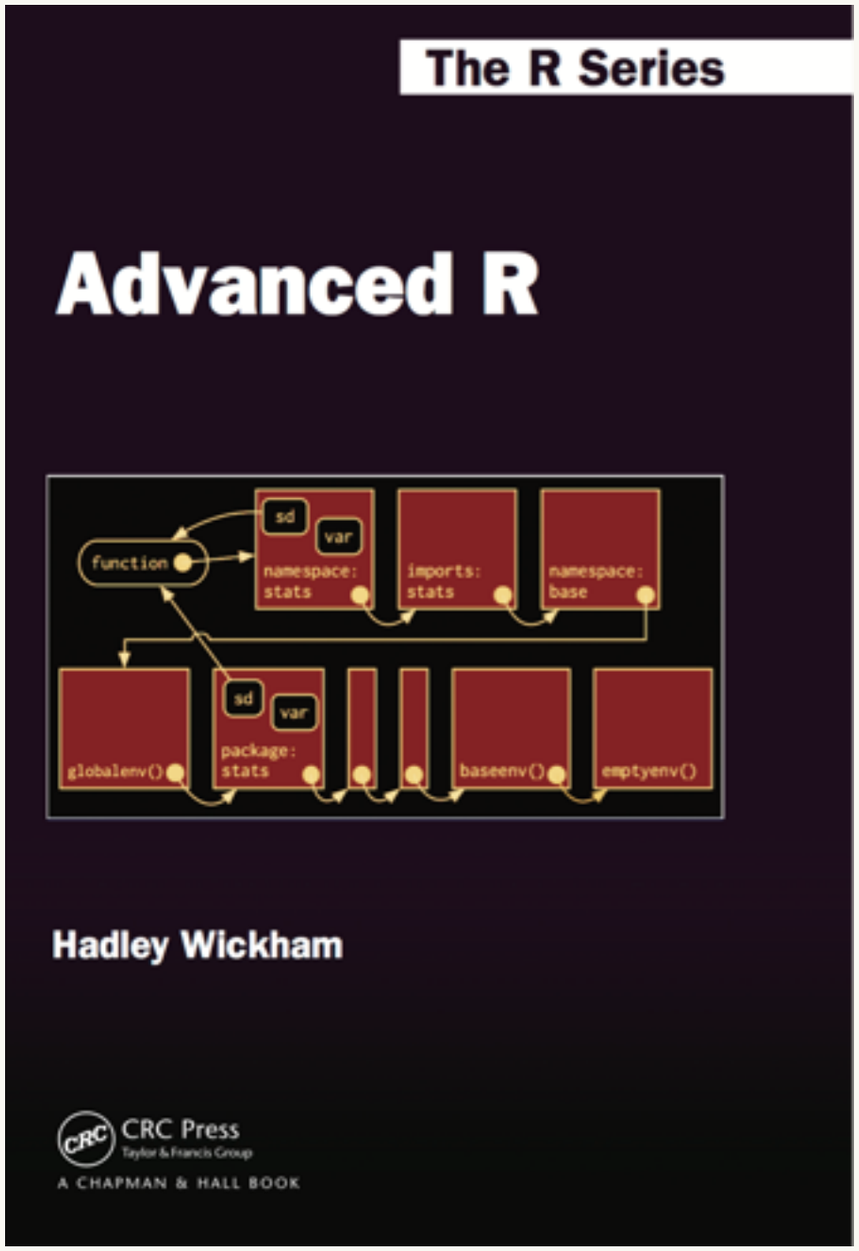

http://adv-r.had.co.nz/

http://amzn.com/1466586966

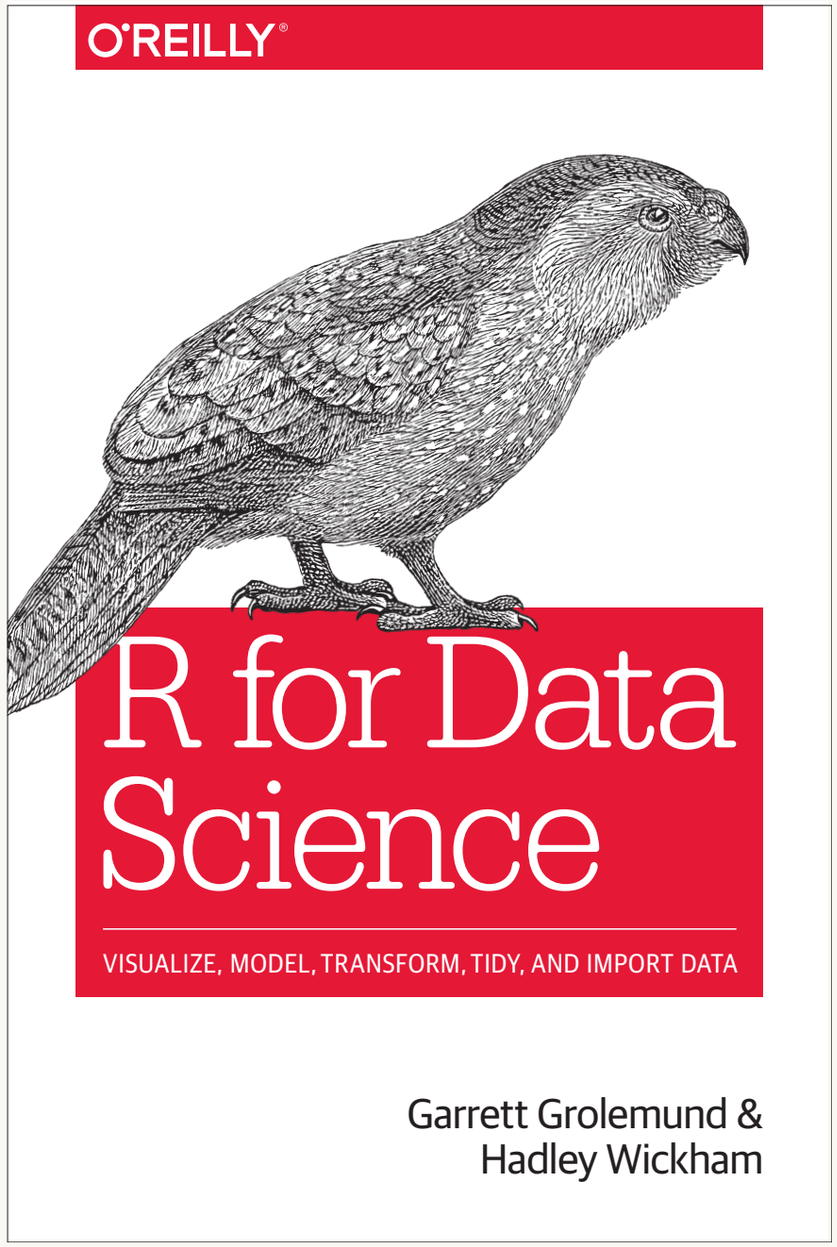

http://r4ds.had.co.nz

https://amzn.com/1491910399

The bad news:
It's going to be frustrating

http://hyperboleandahalf.blogspot.com/2010/09/four-levels-of-social-entrapment.html

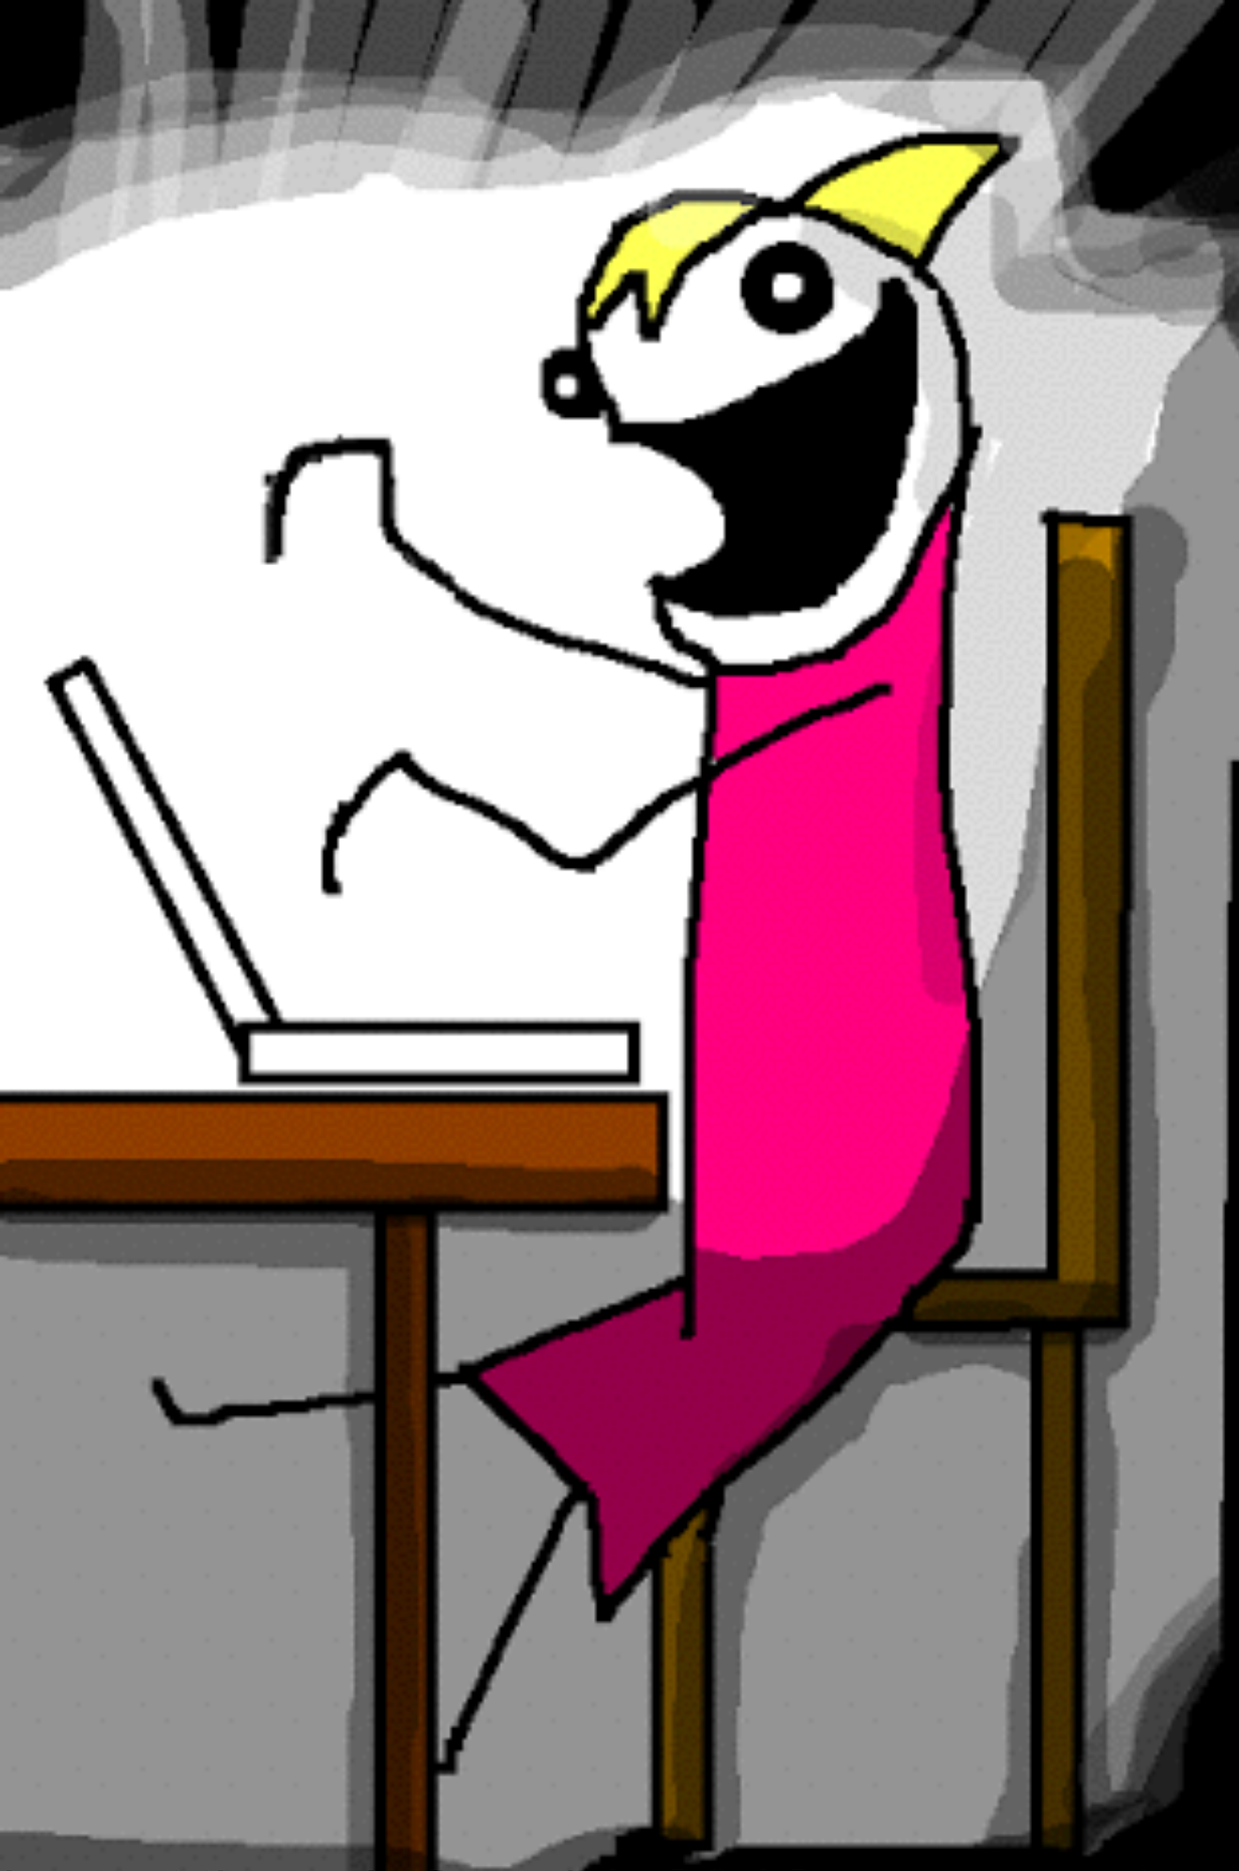

The good news:
Frustration is typical and temporary

http://hyperboleandahalf.blogspot.com/2010/06/this-is-why-ill-never-be-adult.html

Warmups

Don’t expect to know all the answers!

Your turn

What are the four common types of atomic vectors?
(Bonus points for the two uncommon types)

What are the three properties of a vector?

```
typeof(TRUE)
typeof(1L)
typeof(1.5)
typeof("a")

# We'll talk about this (S3) later:
typeof(factor(1:10))
typeof(Sys.Date())
```

Every vector has three properties:

```r
x <- 1:5

# 1. Type:
typeof(x)

# 2. Length
length(x)

# 3. Attributes
attributes(x)
# (we'll come back to those later)
```

~~mode()~~

~~storage.mode()~~

Designed for S compatibility!

Beware built-in is.*() functions

```
# What does is.numeric() do?
x <- 1
y <- 1L
is.integer(x)
is.integer(y)
is.numeric(x)
is.numeric(y)

# What does is.vector() do?
is.vector(1:3)
is.vector(factor(1:3))

# Assume a is.*() does something surprising
```

Purrr provides unsurprising predicates

```
library(purrr)

is_atomic(NULL)
is_vector(factor(1:3))
is_numeric(Sys.Date())

# All purrr predicates are of the form
# function(x) typeof(x) %in% y
```

Coercions

What happens when you try to combine multiple atomic vectors with `c()`? Rank the four types according to which "wins".

What happens when use a logical vector in a numeric context?

Most useful coercion? logical → numeric

```
x <- c(FALSE, TRUE, FALSE, TRUE, TRUE)
sum(x)
mean(x)
```

```
x <- sample(1e4, 1e3)
sum(x > 500)
mean(x > 500)
```

logical ⟶ integer ⟶ numeric ⟶ character

Missing values

What does `NA == NA` return? Why? What should you use instead?

What is `NA_real_`? `NA_integer_`? `NA_character_`?

There isn't a single unknown value

```
age_john <- NA
age_mary <- NA
age_john == age_mary

is.na(x)
!is.na(x) & !is.na(y)

# One NA for each basic atomic vector
typeof(NA)
typeof(NA_real_)
typeof(NA_integer_)
typeof(NA_character_)
```

Your turn

How is a list different from an atomic vector?

How is a data frame different from a matrix?

How do you examine the structure of an object?

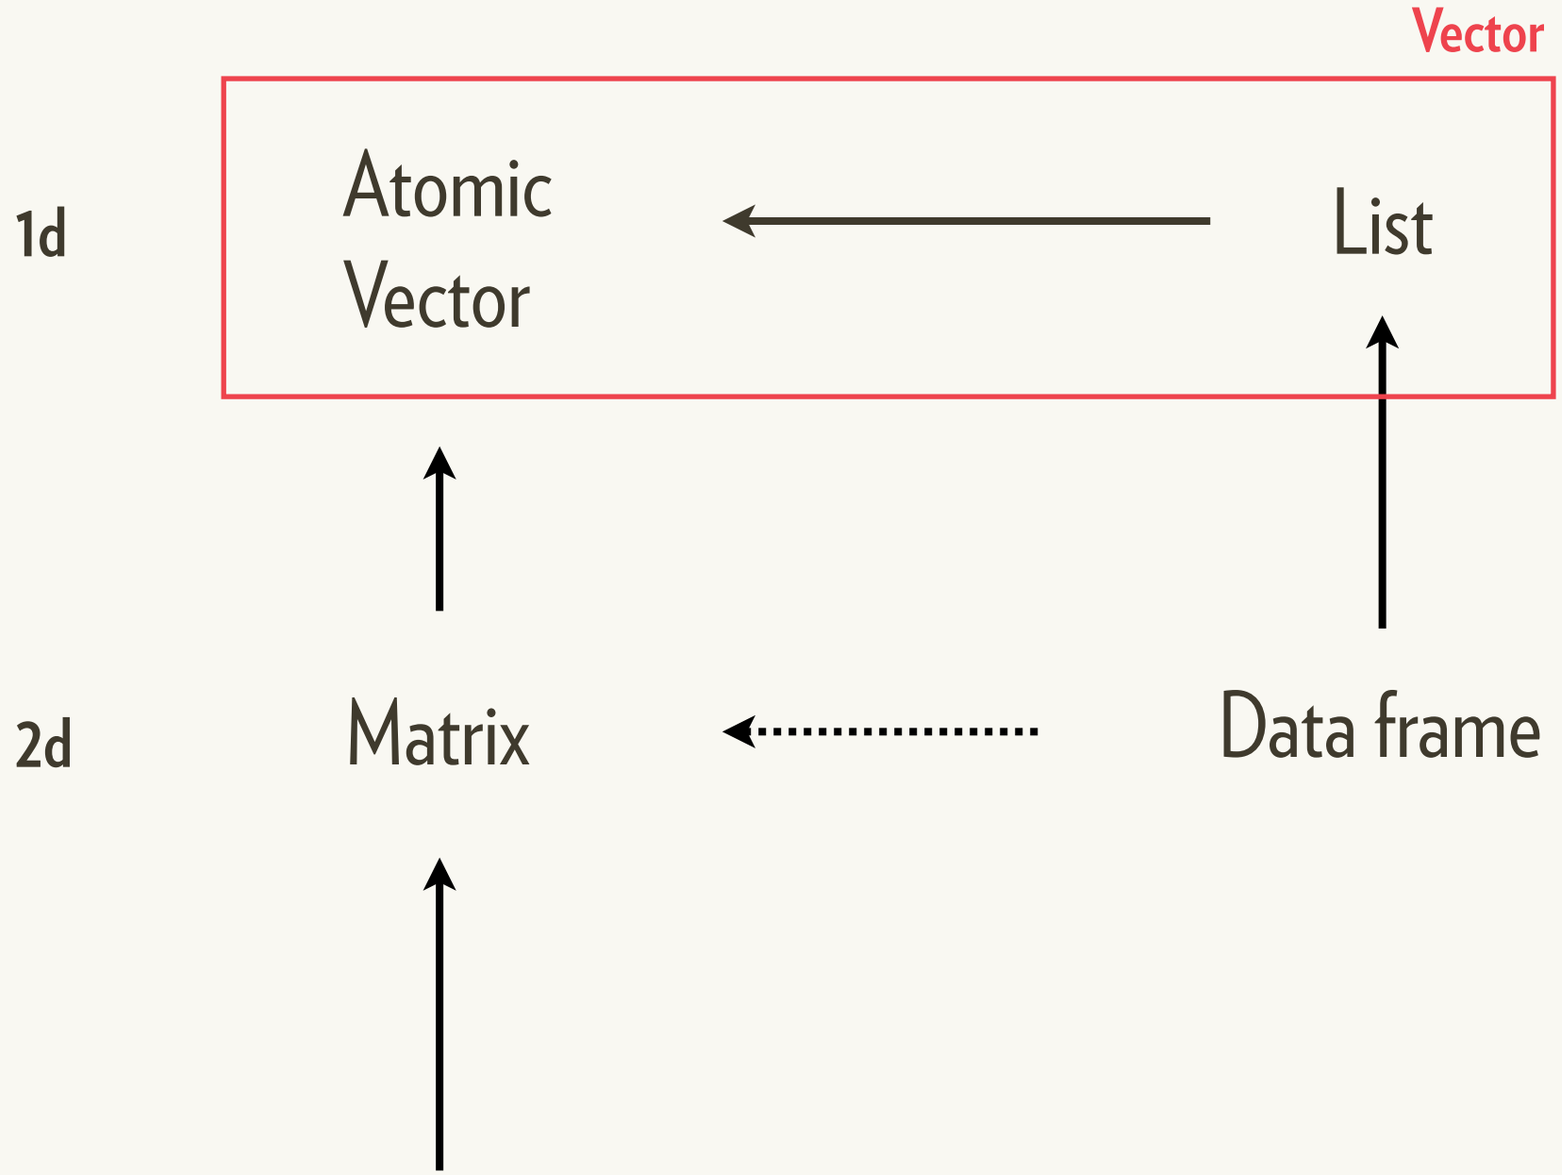

http://adv-r.had.co.nz/Data-structures.html

str()

Your turn

What are the six types of thing that you can put inside []?

What does drop = FALSE do when subsetting a data frame or matrix?

blank	include all
integer	**+ve**: include **0**: drop all **-ve**: exclude
logical	keep TRUEs
character	lookup by name

	Simplifying	Preserving
Vectors	`x[[1]]`	`x[1:4]`
Matrices/ Data frame	`x[, 1]`	`x[, 1, drop = F]`
Lists	`x[[1]]` `x$name`	`x[1]`

Your turn

```
x <- rnorm(1e6)
pryr::object_size(x)
#> 8 MB

# How big will y be?
y <- list(x, x, x)
pryr::object_size(y)

# What if we change a value?
y[[1]][[1]] <- NA
pryr::object_size(y)
```

A name points to an object

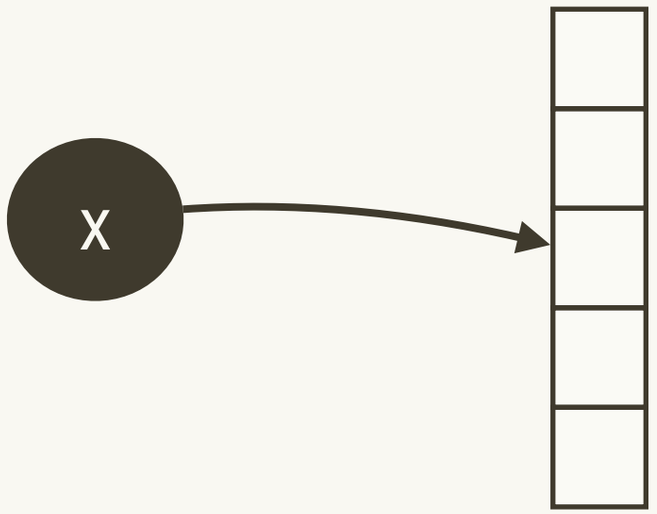

An object can be referenced in multiple places

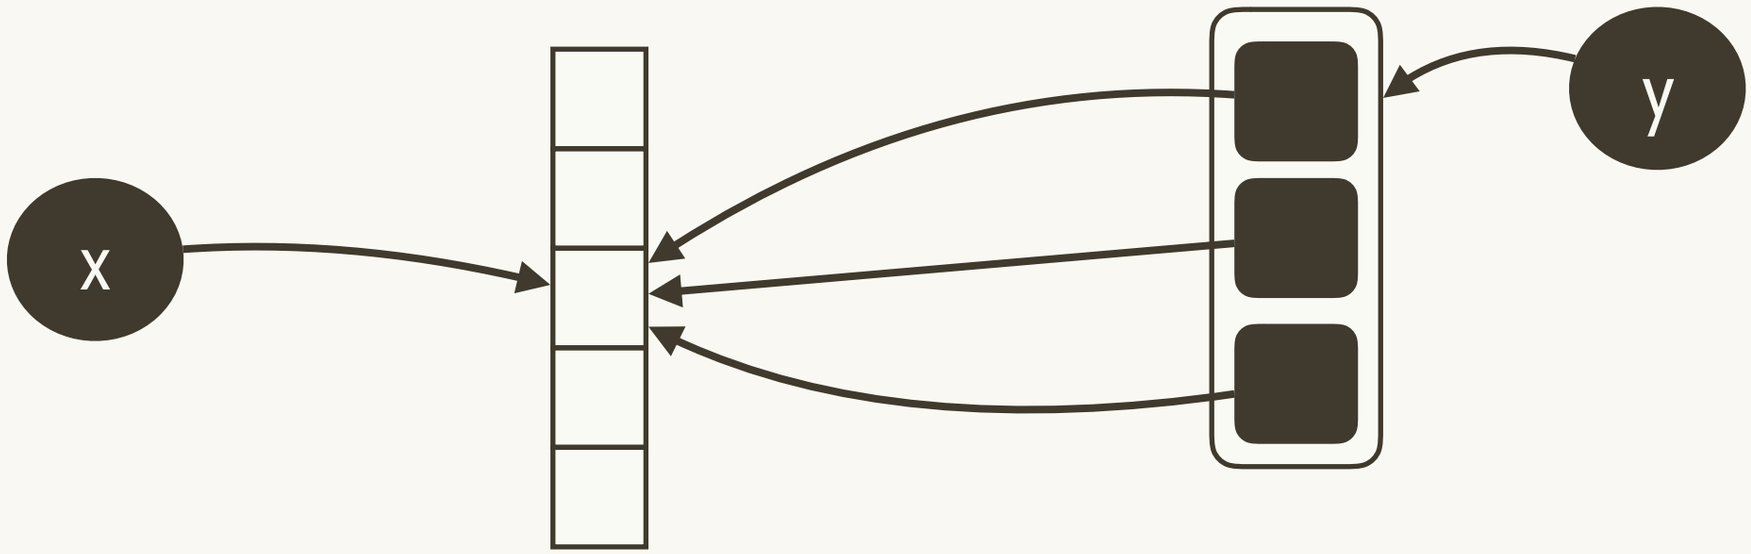

Modifying an object creates a copy

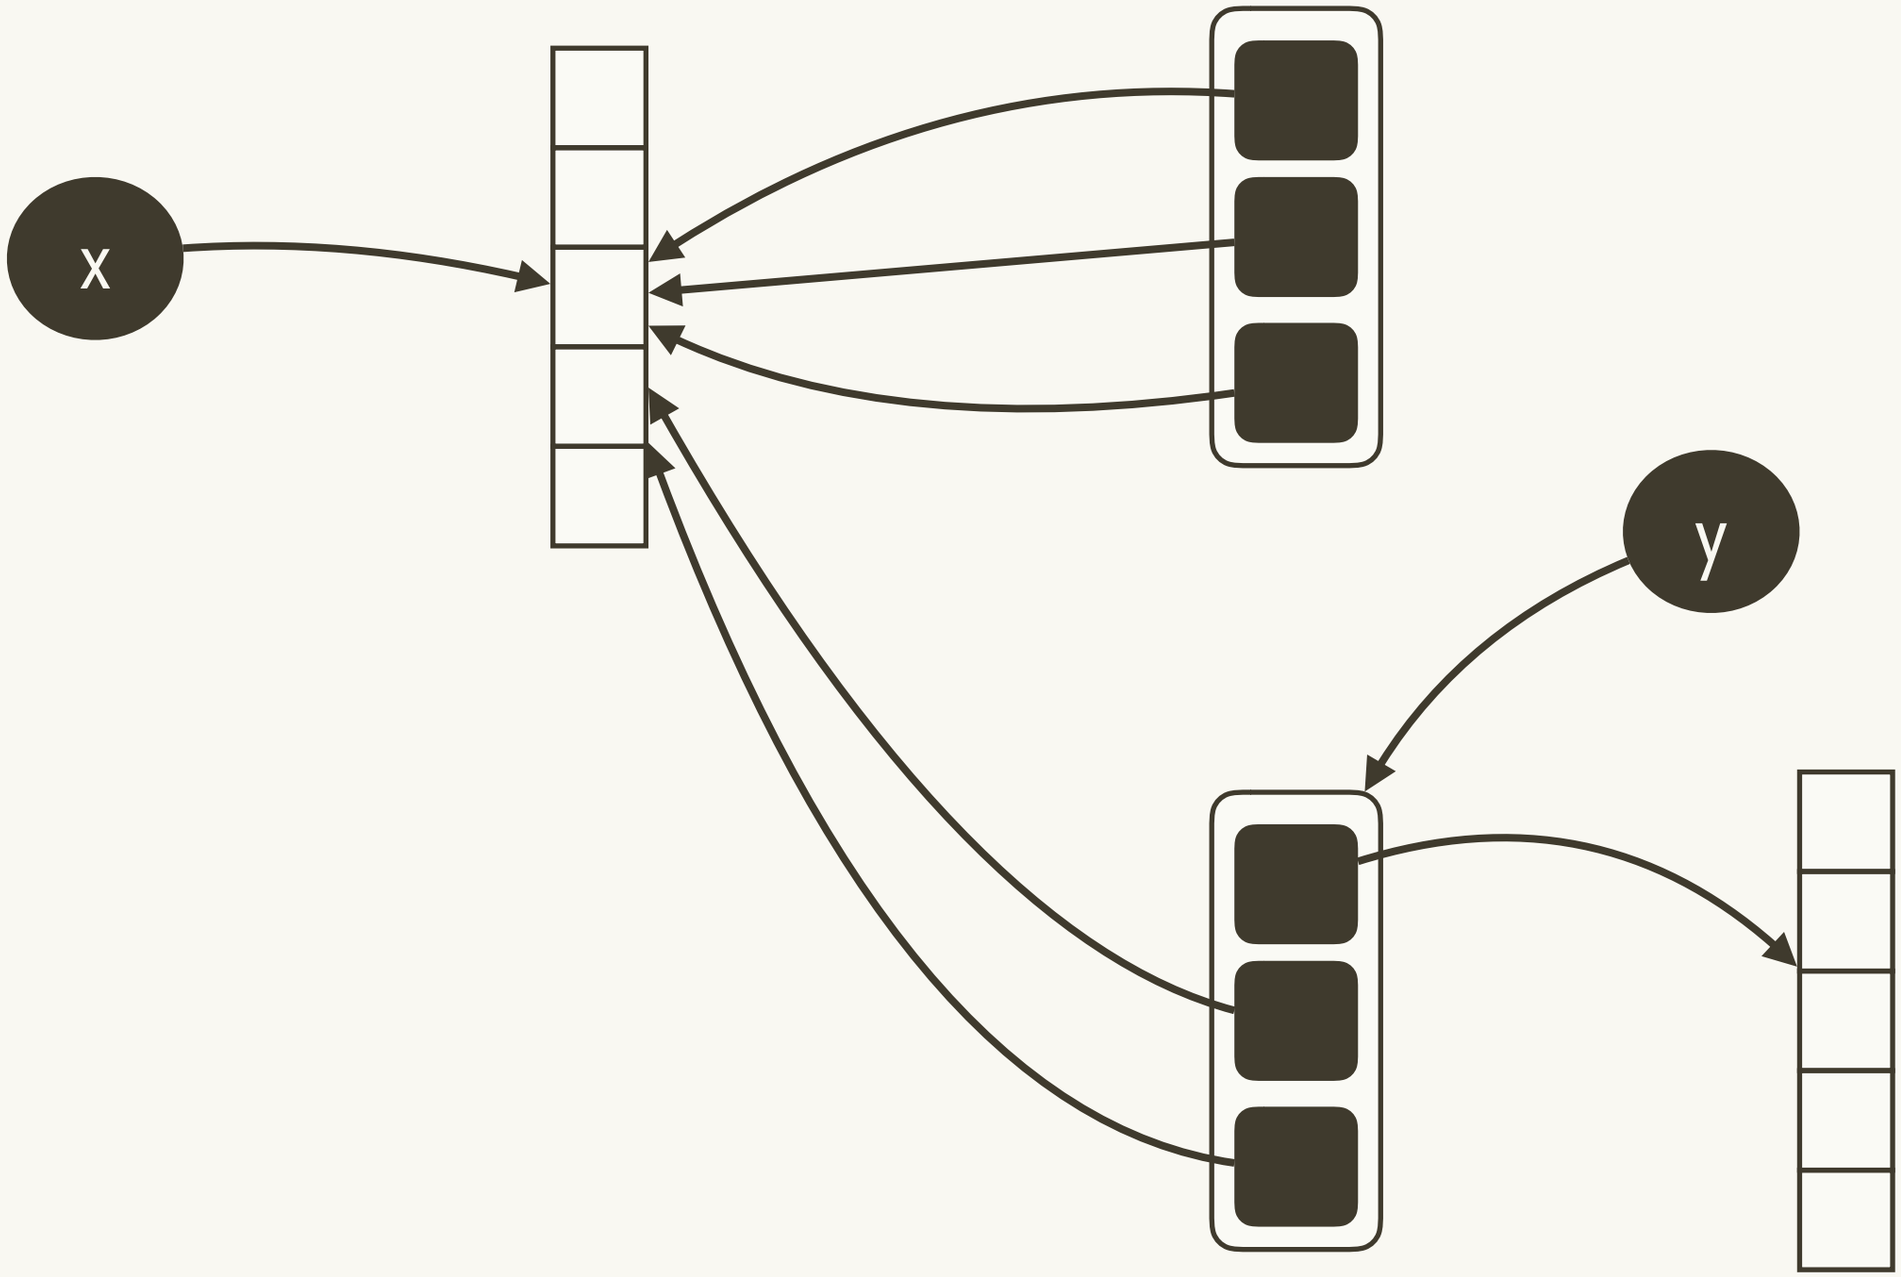

What happens after rm(x) and rm(y)?

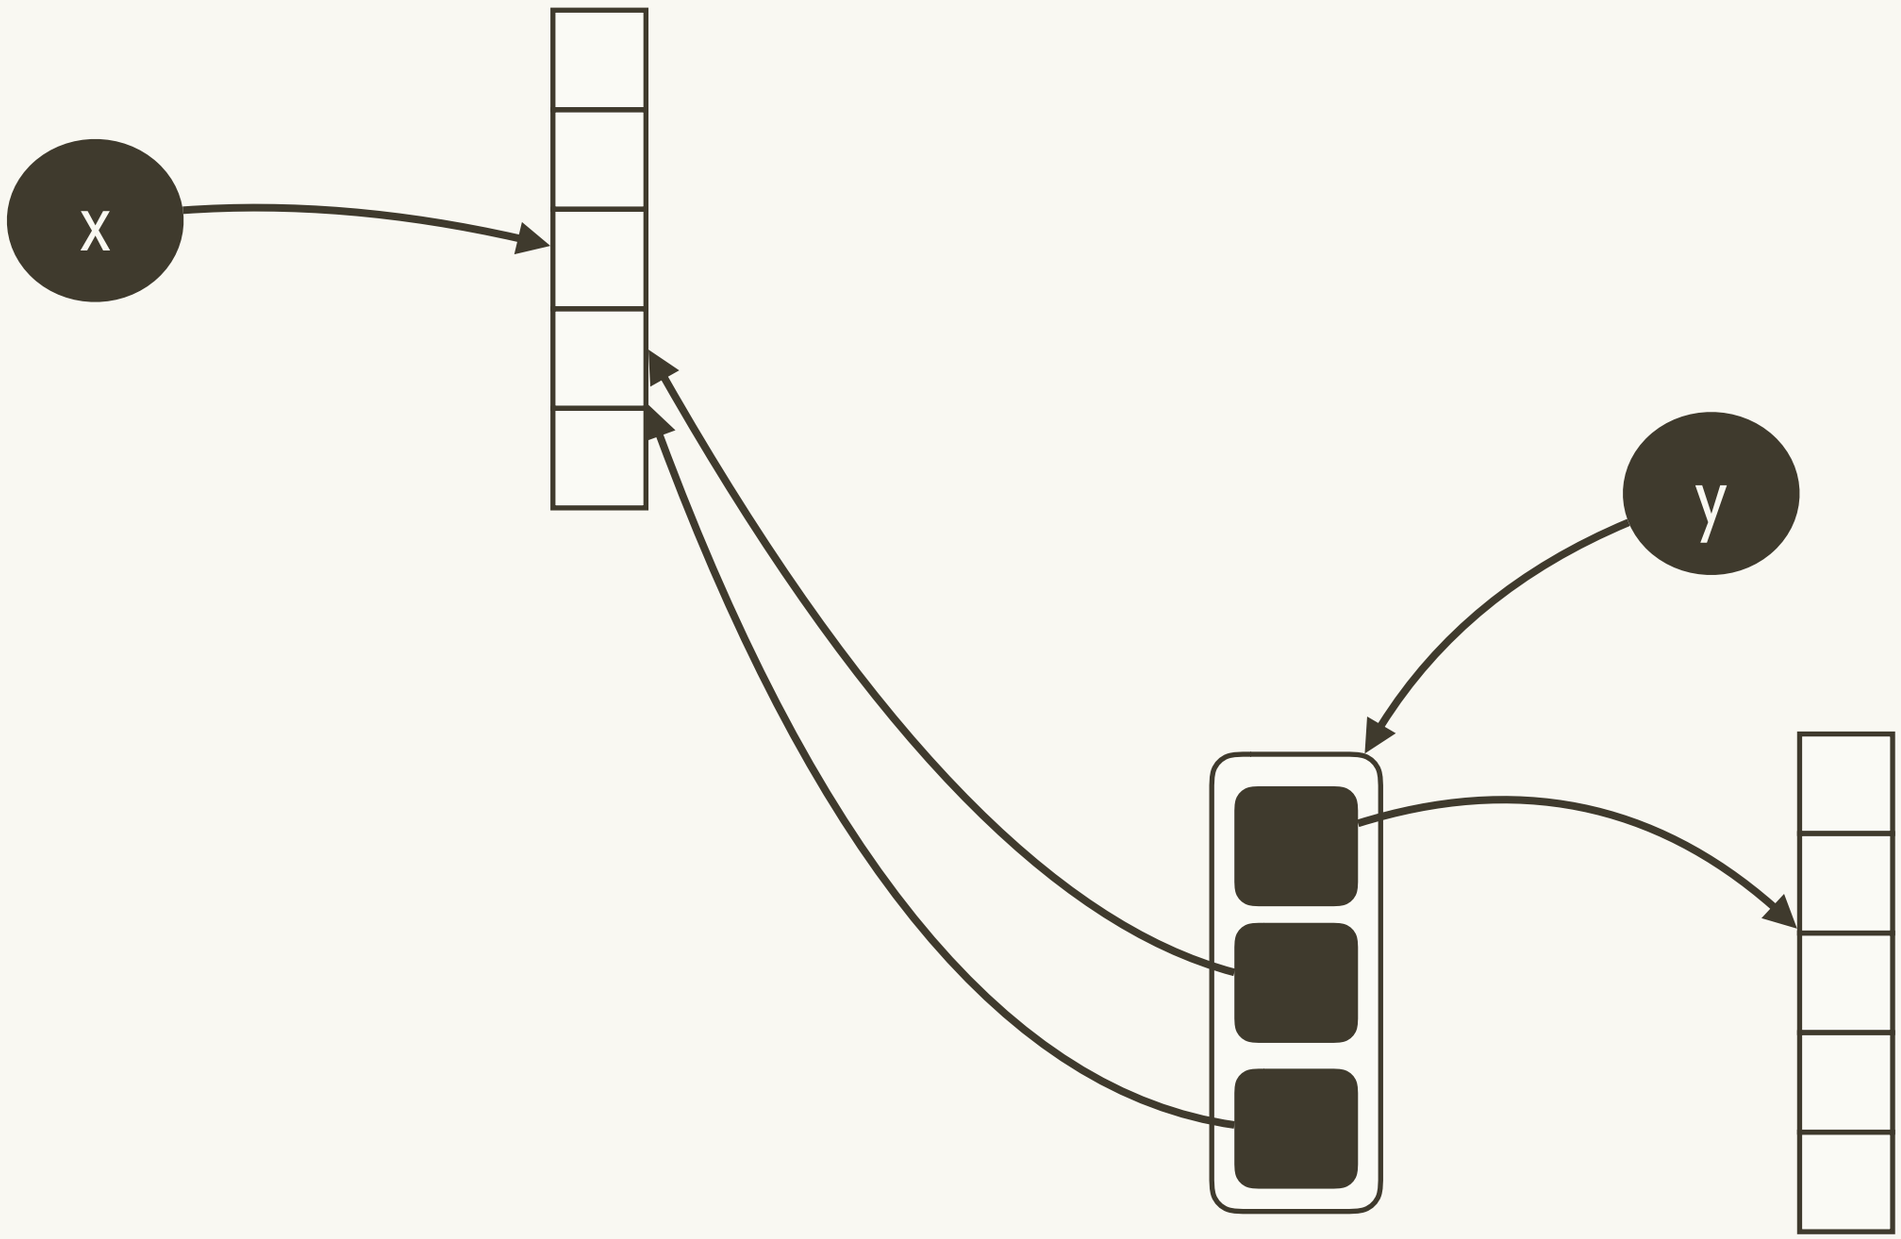

First x and y are removed

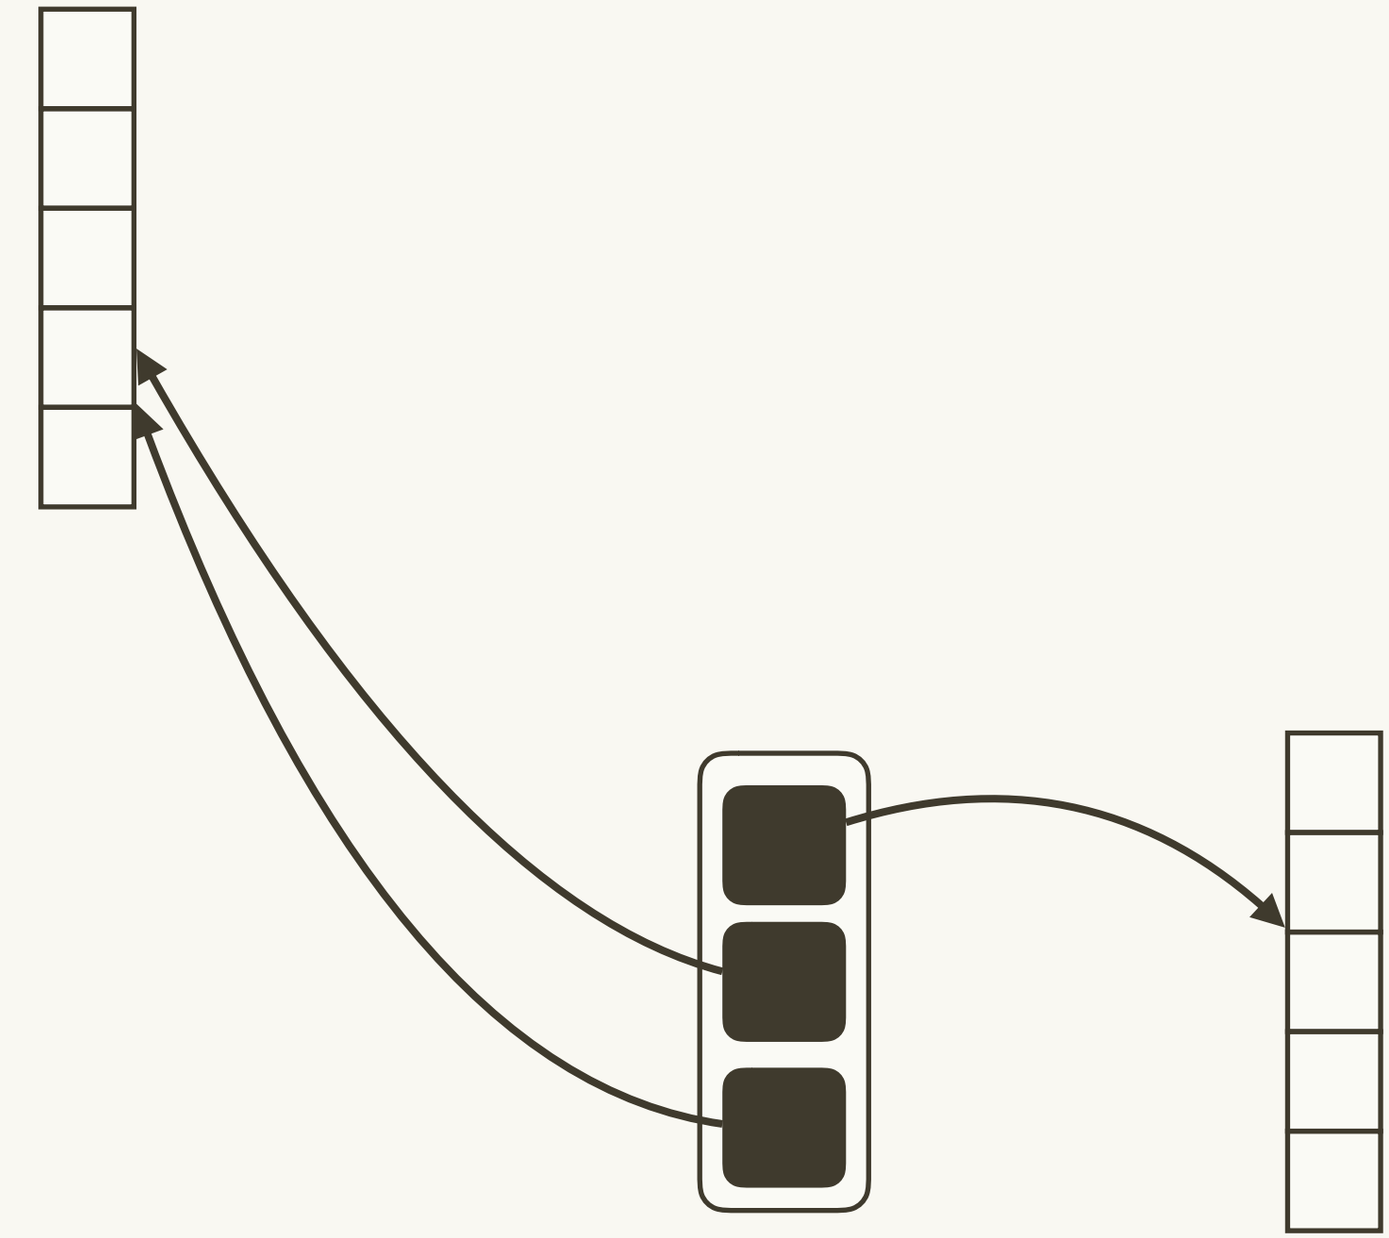

Then the objects they point to

Until nothing is left

The direction of the arrow is important

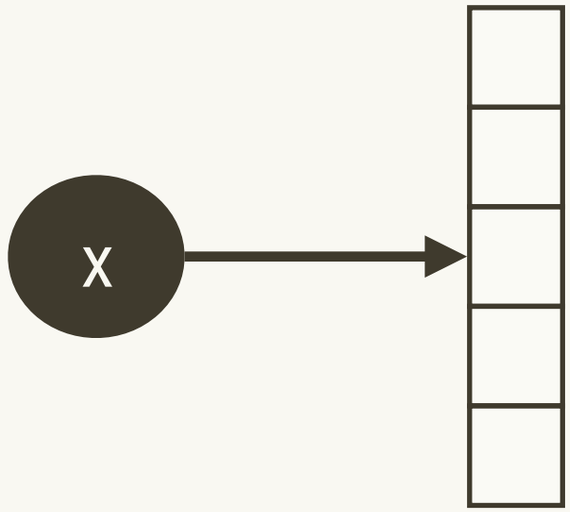

A name “has” an object

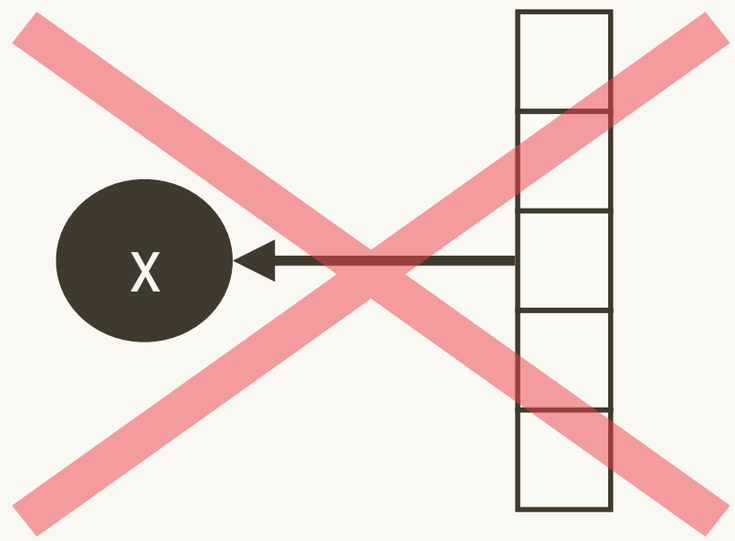

An object doesn’t have a name

RStudio

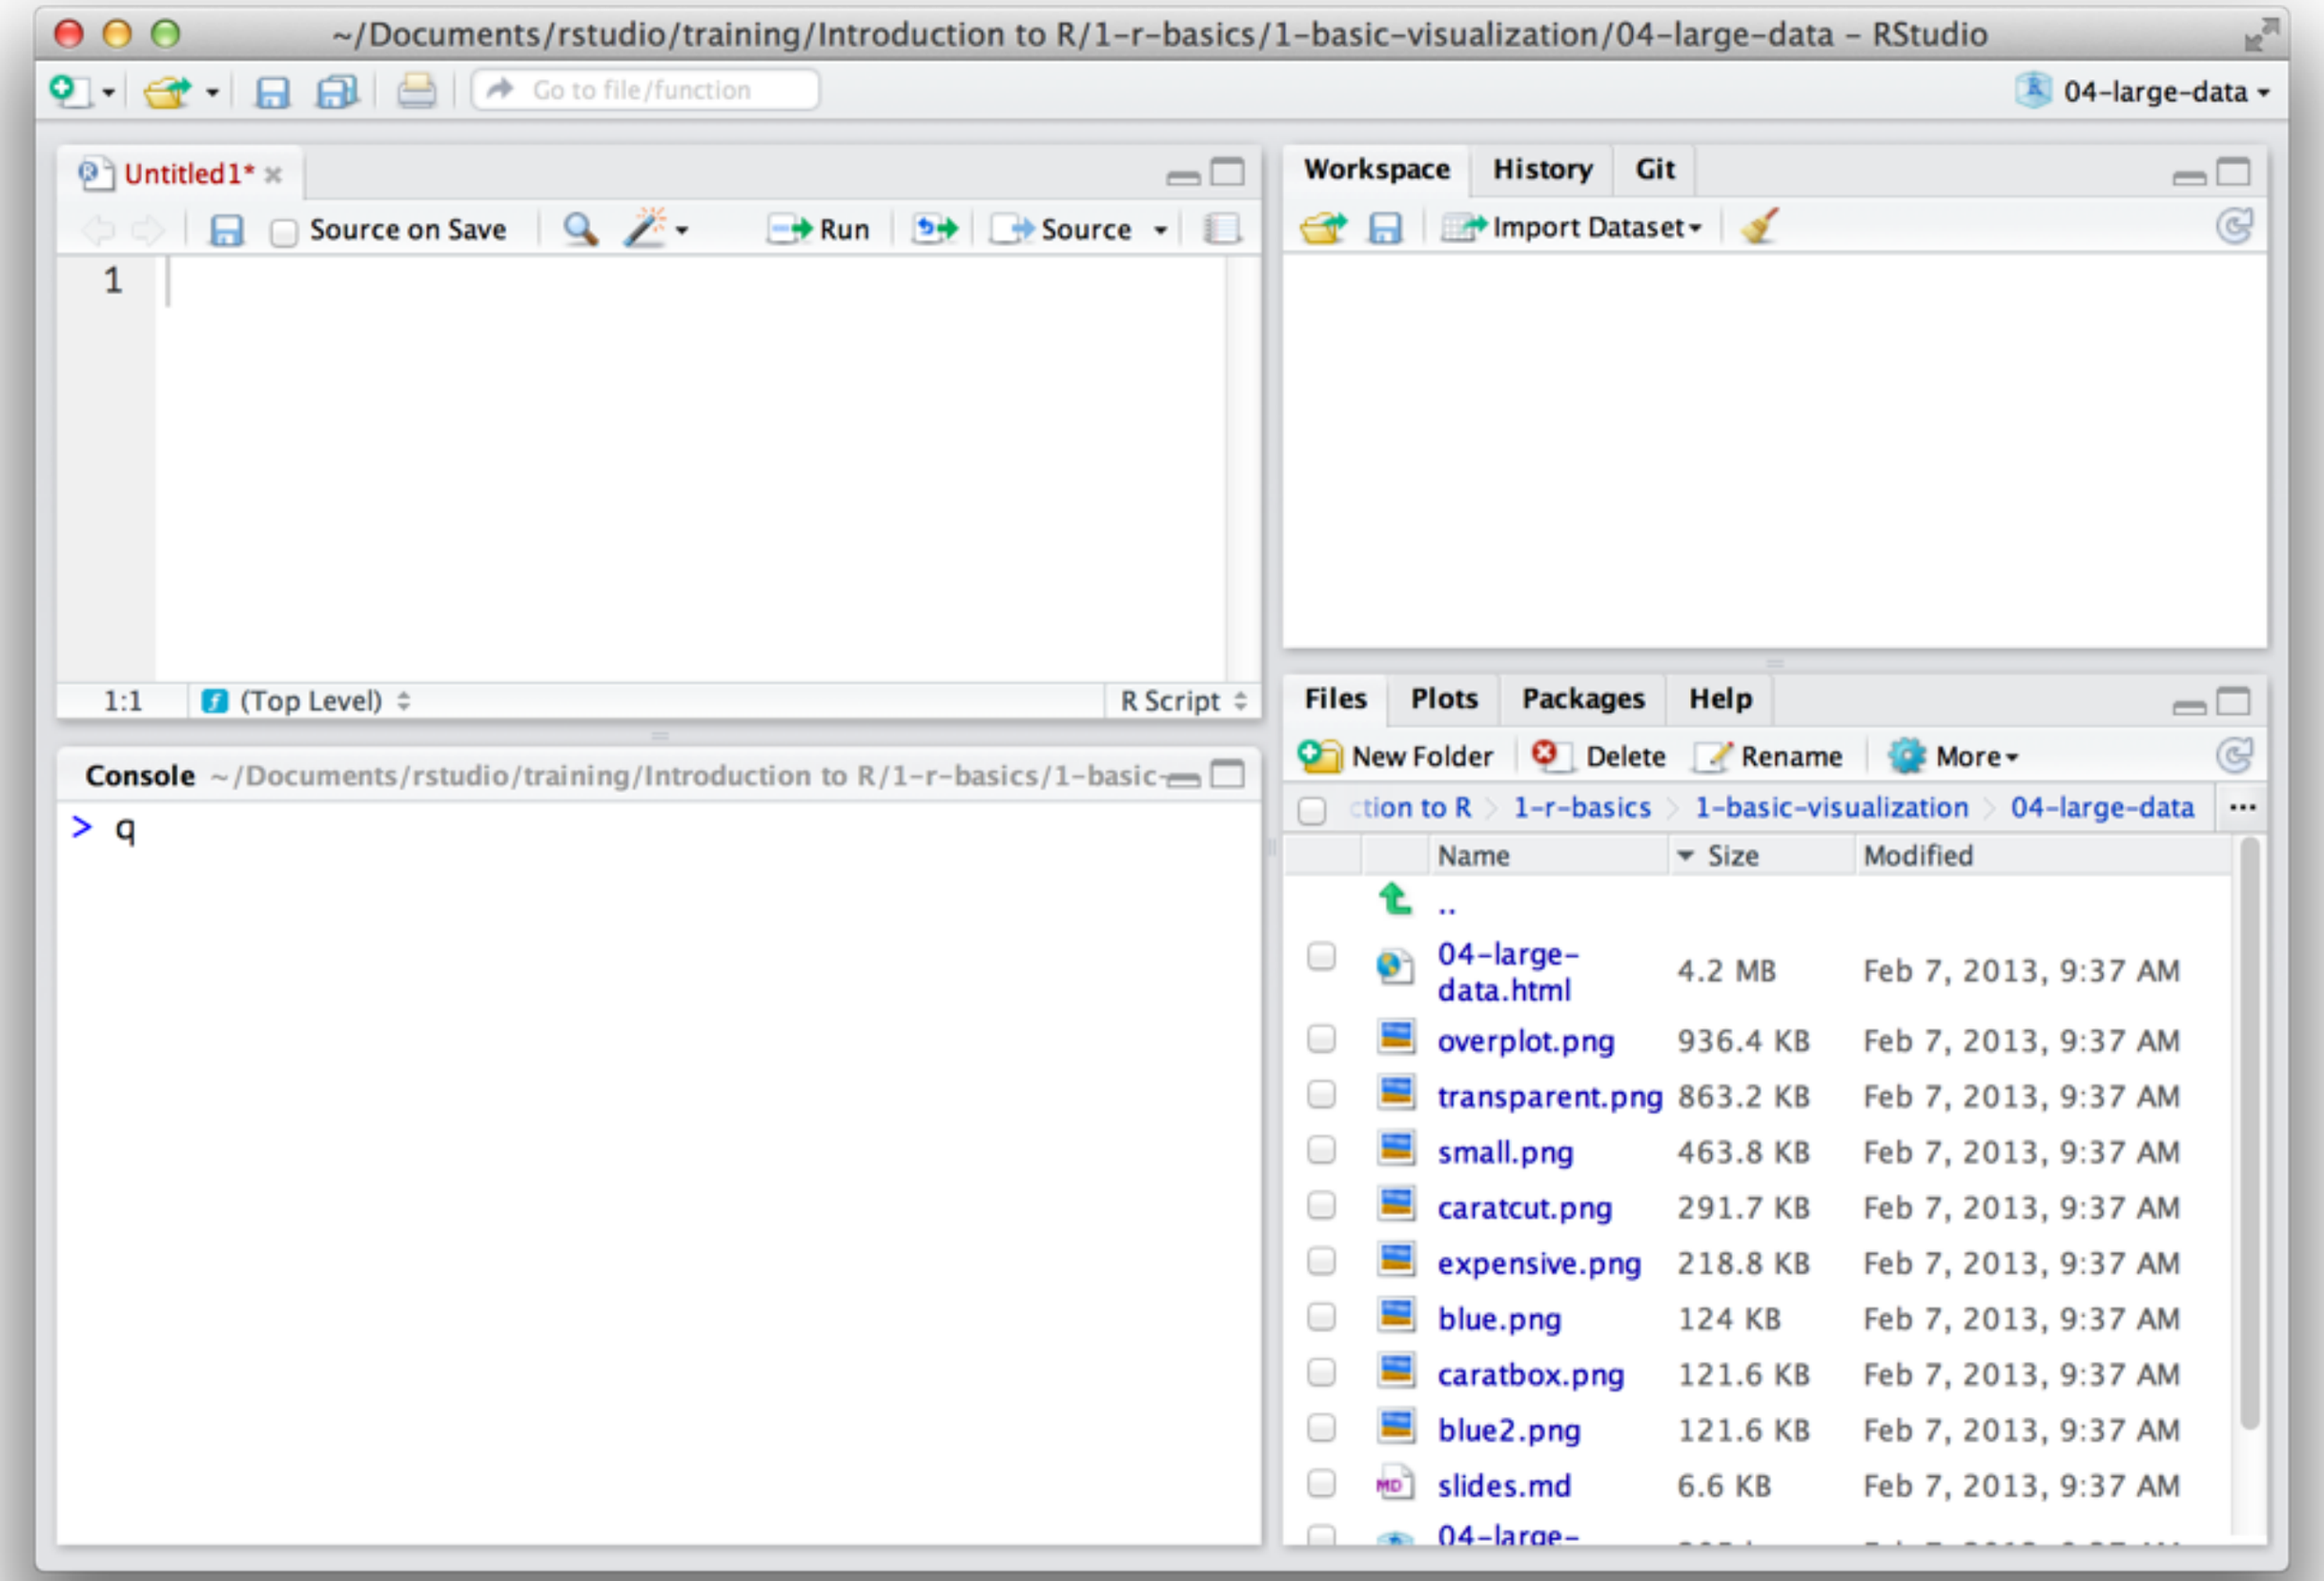
~/Documents/rstudio/training/Introduction to R/1-r-basics/1-basic-visualization/04-large-data - RStudio
Go to file/function
04-large-data
Untitled1*
Source on Save
Run
Source
1
1:1
(Top Level)
R Script
Console ~/Documents/rstudio/training/Introduction to R/1-r-basics/1-basic-
> q
Workspace
History
Git
Import Dataset
Files
Plots
Packages
Help
New Folder
Delete
Rename
More
tion to R
1-r-basics
1-basic-visualization
04-large-data
Name
Size
Modified
..
04-large-data.html 4.2 MB Feb 7, 2013, 9:37 AM
overplot.png 936.4 KB Feb 7, 2013, 9:37 AM
transparent.png 863.2 KB Feb 7, 2013, 9:37 AM
small.png 463.8 KB Feb 7, 2013, 9:37 AM
caratcut.png 291.7 KB Feb 7, 2013, 9:37 AM
expensive.png 218.8 KB Feb 7, 2013, 9:37 AM
blue.png 124 KB Feb 7, 2013, 9:37 AM
caratbox.png 121.6 KB Feb 7, 2013, 9:37 AM
blue2.png 121.6 KB Feb 7, 2013, 9:37 AM
slides.md 6.6 KB Feb 7, 2013, 9:37 AM
04-large-

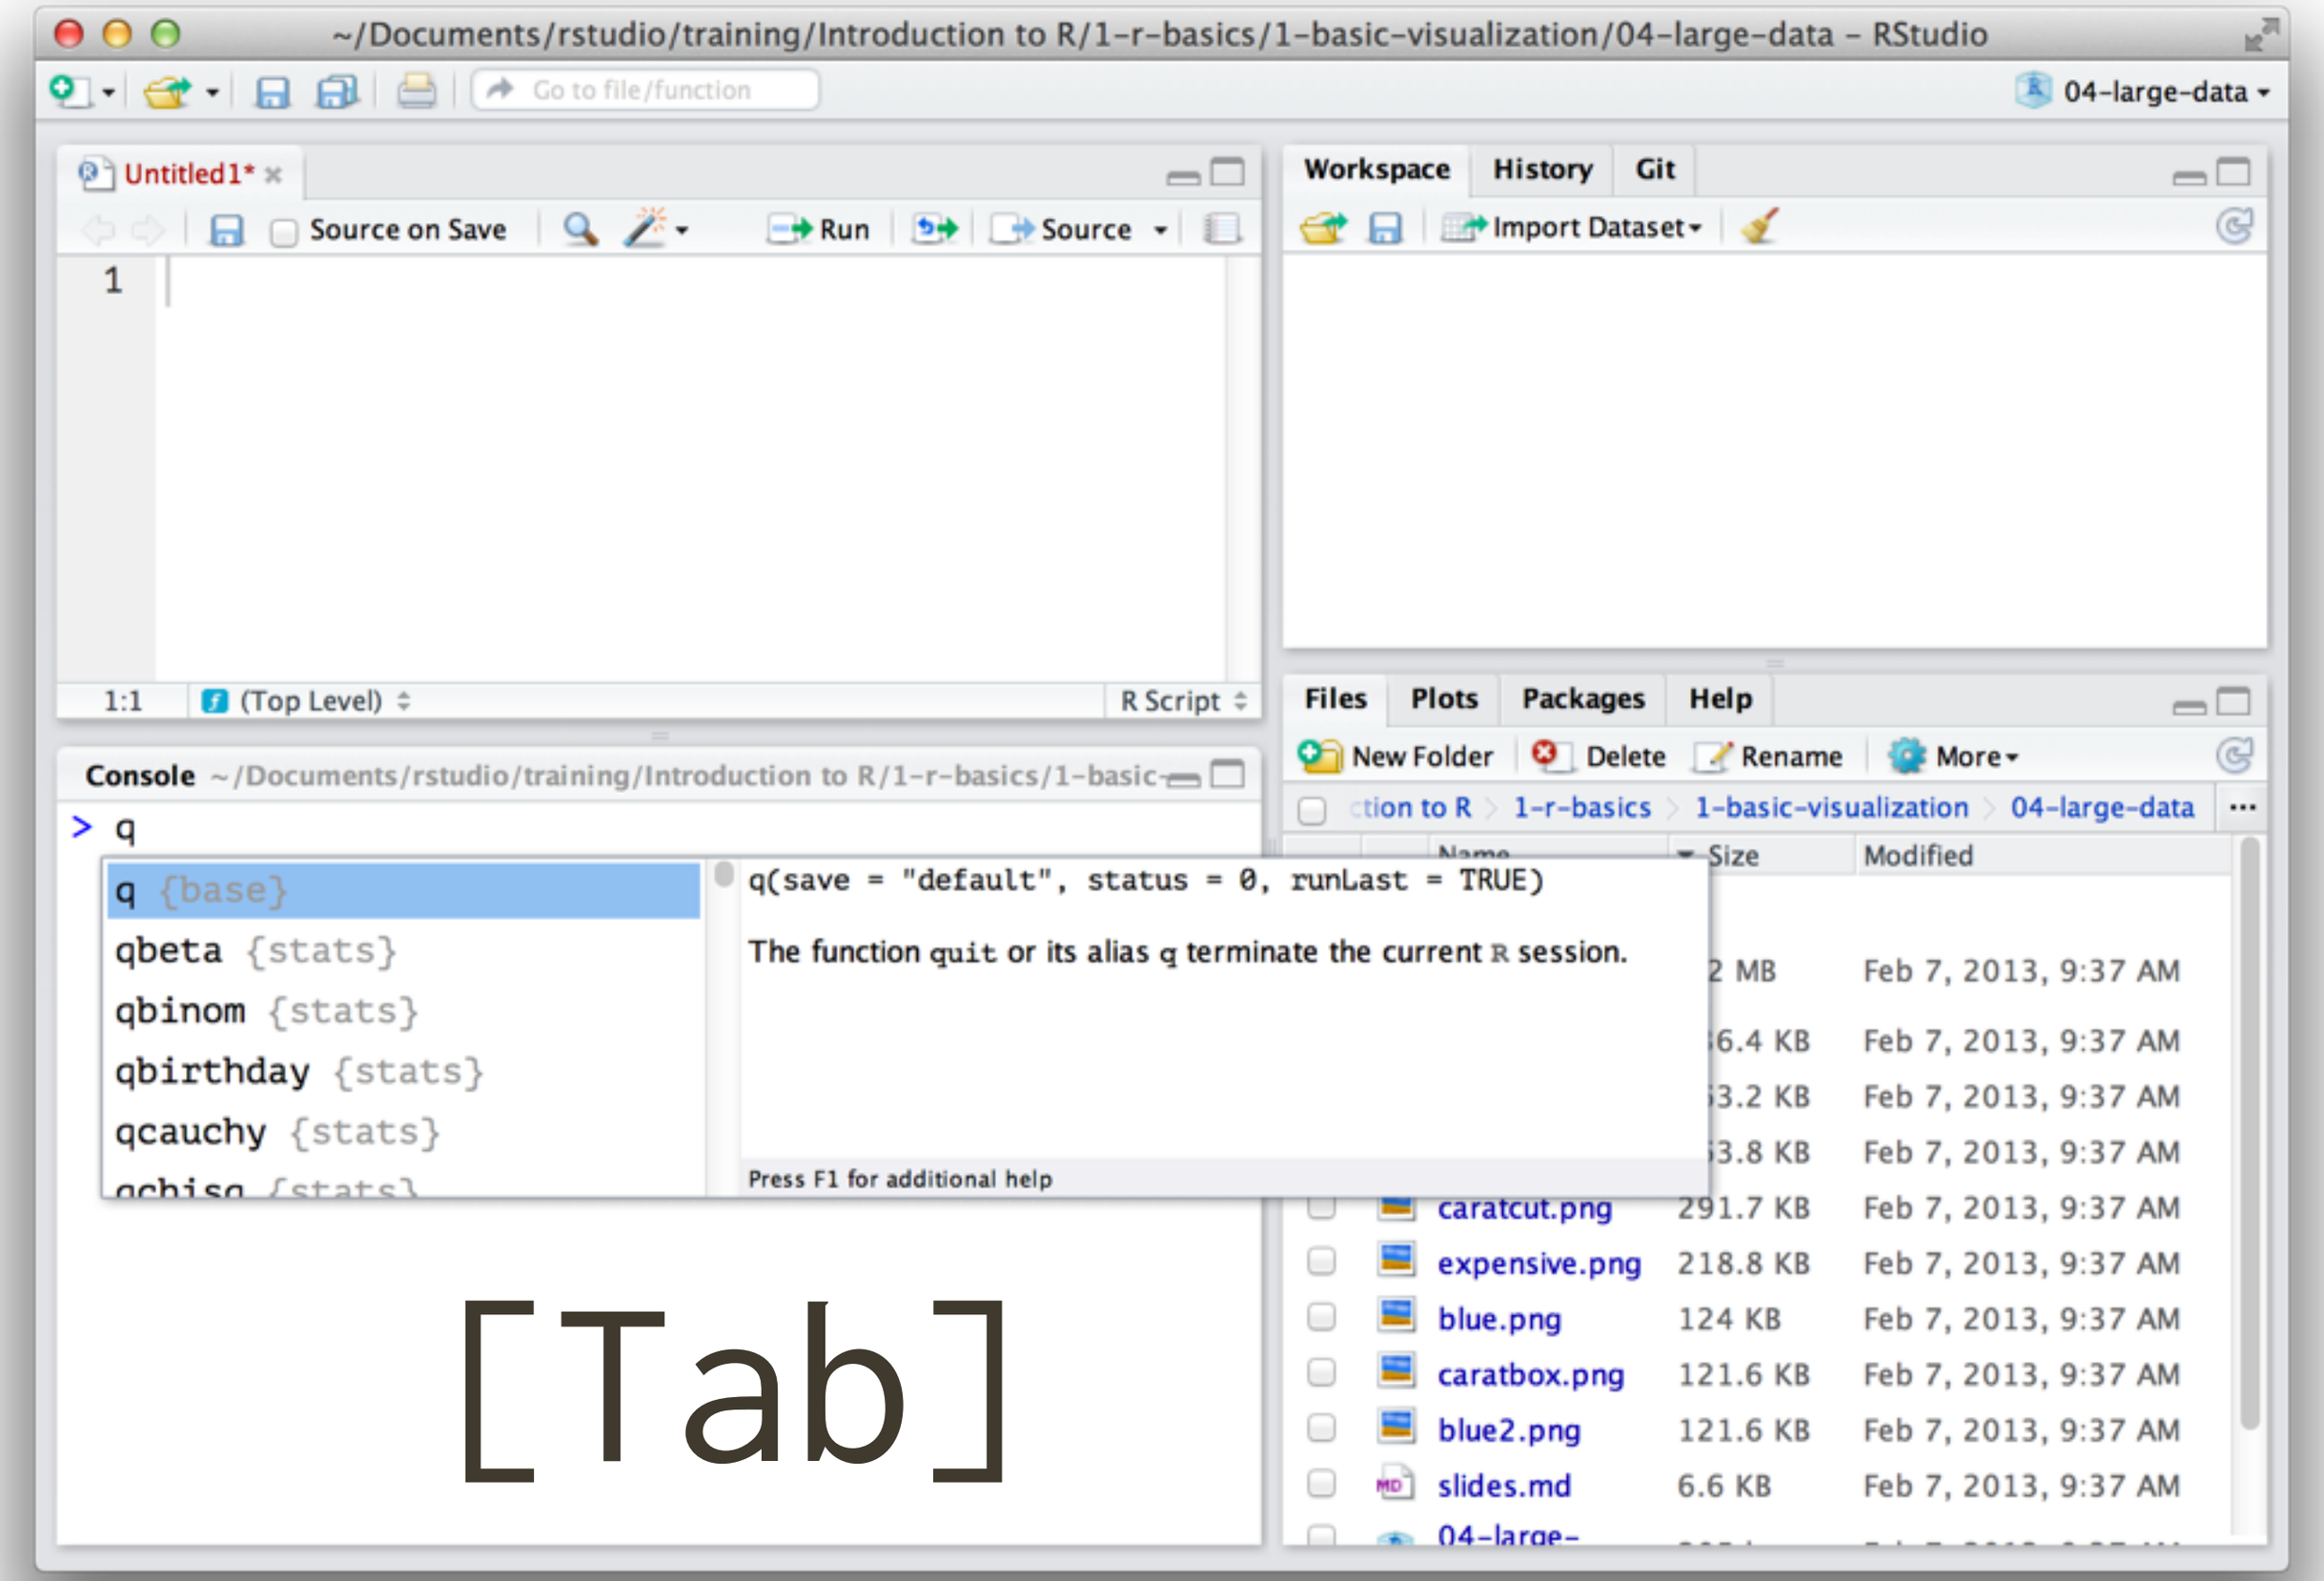
~/Documents/rstudio/training/Introduction to R/1-r-basics/1-basic-visualization/04-large-data - RStudio
Go to file/function
04-large-data
Untitled1*
Source on Save
Run
Source
1
1:1
(Top Level)
R Script
Workspace
History
Git
Import Dataset
Files
Plots
Packages
Help
New Folder
Delete
Rename
More
1-r-basics
1-basic-visualization
04-large-data
Name
Size
Modified
Console ~/Documents/rstudio/training/Introduction to R/1-r-basics/1-basic-
> q
q {base}
qbeta {stats}
qbinom {stats}
qbirthday {stats}
qcauchy {stats}
q(save = "default", status = 0, runLast = TRUE)
The function quit or its alias q terminate the current R session.
Press F1 for additional help
2 MB
Feb 7, 2013, 9:37 AM
6.4 KB
Feb 7, 2013, 9:37 AM
3.2 KB
Feb 7, 2013, 9:37 AM
3.8 KB
Feb 7, 2013, 9:37 AM
caratcut.png
291.7 KB
Feb 7, 2013, 9:37 AM
expensive.png
218.8 KB
Feb 7, 2013, 9:37 AM
blue.png
124 KB
Feb 7, 2013, 9:37 AM
caratbox.png
121.6 KB
Feb 7, 2013, 9:37 AM
blue2.png
121.6 KB
Feb 7, 2013, 9:37 AM
slides.md
6.6 KB
Feb 7, 2013, 9:37 AM
[Tab]

~/Documents/rstudio/training/Introduction to R/1-r-basics/1-basic-visualization/04-large-data - RStudio

Go to file/function

04-large-data

Untitled1*

Source on Save | Run | Source

```
qplot(table - depth, data = diamonds,
qplot(day, data = email)
qplot(day, mails, data = daily, geom = "line", colo
qplot(day, mails, data = daily, geom = "smooth", co
qplot(day, variants, data = daily, geom = "line", c
qplot(wday, hour, data = wh, size = freq)
qplot(mpg, wt, data = mtcars)
qplot(mpg, wt, data = mtcars, colour = cyl)
```

```
> q
```

Workspace | History | Git

Import Dataset

Files | Plots | Packages | Help

New Folder | Delete | Rename | More

ction to R > 1-r-basics > 1-basic-visualization > 04-large-data

	Name	Size	Modified
	..		
	04-large-data.html	4.2 MB	Feb 7, 2013, 9:37 AM
	overplot.png	936.4 KB	Feb 7, 2013, 9:37 AM
	transparent.png	863.2 KB	Feb 7, 2013, 9:37 AM
	small.png	463.8 KB	Feb 7, 2013, 9:37 AM
	caratcut.png	291.7 KB	Feb 7, 2013, 9:37 AM
		:B	Feb 7, 2013, 9:37 AM
			Feb 7, 2013, 9:37 AM
		:B	Feb 7, 2013, 9:37 AM
		:B	Feb 7, 2013, 9:37 AM
			Feb 7, 2013, 9:37 AM
	04-large-		

[Cmd/Ctrl + ↑]

~/Documents/rstudio/training/Introduction to R/1-r-basics/1-basic-visualization/04-large-data - RStudio

Go to file/function

04-large-data

Untitled1*

Source on Save

Run

Source

```
1  library(ggplot2)
```

1:17 (Top Level) R Script

Console ~/Documents/rstudio/training/Introduction to R/1-r-basics/1-basic-

```
>
```

Workspace History Git

Import Dataset

Files Plots Packages Help

New Folder Delete Rename More

tion to R 1-r-basics 1-basic-visualization 04-large-data

	Name	Size	Modified
	..		
	04-large-data.html	4.2 MB	Feb 7, 2013, 9:37 AM
	overplot.png	936.4 KB	Feb 7, 2013, 9:37 AM
	transparent.png	863.2 KB	Feb 7, 2013, 9:37 AM
	small.png	463.8 KB	Feb 7, 2013, 9:37 AM
	caratcut.png	291.7 KB	Feb 7, 2013, 9:37 AM
	expensive.png	218.8 KB	Feb 7, 2013, 9:37 AM
	blue.png	124 KB	Feb 7, 2013, 9:37 AM
	caratbox.png	121.6 KB	Feb 7, 2013, 9:37 AM
	blue2.png	121.6 KB	Feb 7, 2013, 9:37 AM
	slides.md	6.6 KB	Feb 7, 2013, 9:37 AM
	04-large-		

~/Documents/rstudio/training/Introduction to R/1-r-basics/1-basic-visualization/04-large-data - RStudio

Go to file/function

04-large-data

Untitled1*

Source on Save | Run | Source

```
1  library(ggplot2)
```

[Cmd/ctrl + enter]

1:17 (Top Level) R Script

Console ~/Documents/rstudio/training/Introduction to R/1-r-basics/1-basic-

```
> library(ggplot2)
>
```

Workspace | History | Git

Import Dataset

Files | Plots | Packages | Help

New Folder | Delete | Rename | More

ction to R › 1-r-basics › 1-basic-visualization › 04-large-data

Name	Size	Modified
..		
04-large-data.html	4.2 MB	Feb 7, 2013, 9:37 AM
overplot.png	936.4 KB	Feb 7, 2013, 9:37 AM
transparent.png	863.2 KB	Feb 7, 2013, 9:37 AM
small.png	463.8 KB	Feb 7, 2013, 9:37 AM
caratcut.png	291.7 KB	Feb 7, 2013, 9:37 AM
expensive.png	218.8 KB	Feb 7, 2013, 9:37 AM
blue.png	124 KB	Feb 7, 2013, 9:37 AM
caratbox.png	121.6 KB	Feb 7, 2013, 9:37 AM
blue2.png	121.6 KB	Feb 7, 2013, 9:37 AM
slides.md	6.6 KB	Feb 7, 2013, 9:37 AM
04-large-		

~/Documents/rstudio/training/Introduction to R/1-r-basics/1-basic-visualization/04-large-data - RStudio

Go to file/function

04-large-data

Untitled1*

Source on Save | Run | Source

```
1  library(ggplot2)
```

[Cmd/ctrl + shift + F10]

1:17 (Top Level) R Script

Console ~/Documents/rstudio/training/Introduction to R/1-r-basics/1-basic-

```
> library(ggplot2)
>

Restarting R session...

>
```

Workspace | History | Git

Import Dataset

Files | Plots | Packages | Help

New Folder | Delete | Rename | More

tion to R › 1-r-basics › 1-basic-visualization › 04-large-data

Name	Size	Modified
..		
04-large-data.html	4.2 MB	Feb 7, 2013, 9:37 AM
overplot.png	936.4 KB	Feb 7, 2013, 9:37 AM
transparent.png	863.2 KB	Feb 7, 2013, 9:37 AM
small.png	463.8 KB	Feb 7, 2013, 9:37 AM
caratcut.png	291.7 KB	Feb 7, 2013, 9:37 AM
expensive.png	218.8 KB	Feb 7, 2013, 9:37 AM
blue.png	124 KB	Feb 7, 2013, 9:37 AM
caratbox.png	121.6 KB	Feb 7, 2013, 9:37 AM
blue2.png	121.6 KB	Feb 7, 2013, 9:37 AM
slides.md	6.6 KB	Feb 7, 2013, 9:37 AM
04-large-		

Options

General

Code

Appearance

Pane Layout

Packages

Sweave

Spelling

Git/SVN

Publishing

Default working directory (when not in a project):

~ Browse...

- [x] Restore most recently opened project at startup
- [x] Restore previously open source documents at startup
- [ ] Restore .RData into workspace at startup

Save workspace to .RData on exit: Never

- [x] Always save history (even when not saving .RData)
- [x] Remove duplicate entries in history
- [ ] Use debug error handler only when my code contains errors
- [ ] Automatically expand tracebacks in error inspector

Default text encoding:

UTF-8 Change...

- [x] Automatically notify me of updates to RStudio

OK Cancel Apply

~/Documents/rstudio/training/Introduction to R/1-r-basics/1-basic-visualization/04-large-data - RStudio

Go to file/function

04-large-data

Untitled1*

Source on Save | Run | Source

```
1  library(ggplot2)
```

[Alt + shift + K]

1:17 (Top Level) R Script

Workspace | History | Git

Import Dataset

Files | Plots | Packages | Help

New Folder | Delete | Rename | More

tion to R › 1-r-basics › 1-basic-visualization › 04-large-data

Console ~/Documents/rstudio/training/Introduction to R/1-r-basics/1-basic-

```
> library(ggplot2)
>
Restarting R session...

>
```

Name	Size	Modified
..		
04-large-data.html	4.2 MB	Feb 7, 2013, 9:37 AM
overplot.png	936.4 KB	Feb 7, 2013, 9:37 AM
transparent.png	863.2 KB	Feb 7, 2013, 9:37 AM
small.png	463.8 KB	Feb 7, 2013, 9:37 AM
caratcut.png	291.7 KB	Feb 7, 2013, 9:37 AM
expensive.png	218.8 KB	Feb 7, 2013, 9:37 AM
blue.png	124 KB	Feb 7, 2013, 9:37 AM
caratbox.png	121.6 KB	Feb 7, 2013, 9:37 AM
blue2.png	121.6 KB	Feb 7, 2013, 9:37 AM
slides.md	6.6 KB	Feb 7, 2013, 9:37 AM
04-large-		

Your turn

Check that you're familiar with the basic workflow outlined in http://r4ds.had.co.nz/workflow-scripts.html

Skim https://twitter.com/rstudiotips – which one do you think is most useful?